

Bl~~ack~~jack AI

Ryan Karch (karchr), Dominic Beyer (beyerd), Angelica Loshak (lossha), Michael Dong (dongm2)

Fall 2023

1 Project Description

1.1 Problem Solved

Blackjack AI gives users a place where they can test their blackjack skills alongside artificial intelligence agents, including being able to upload their own artificial intelligence agents. The main objective of this project is to provide an entertaining platform that encourages skill development in blackjack. It enhances the strategic thinking of players by giving them a sense of competition from a challenging AI opponent. The website aims to create a user-friendly environment where players can interact with AI and gain experience with betting cryptocurrency in a safe and secure space. The project also addresses the different skill levels of players by allowing them to adjust the intelligence levels of the bot that they are playing with. This ensures that players from varying blackjack backgrounds can benefit and grow from Blackjack AI.

1.2 Background

Blackjack is one of the most popular card games in casinos. However as online casinos become a significant source of revenue for the gambling industry, more players are falling to gambling addiction. The accessibility and convenience that come with online gambling make it easier than ever before for players to make impulsive decisions and run into financial problems. Along with that, deception among online gambling platforms contributes to players falling victim to underage gambling, scams, and fraud.

Blackjack AI aims to bring users an alternative way to play blackjack with friends and AI players without losing large amounts of money or risking their safety and privacy. We prioritize safety and responsibility by having users bet in Wei instead of full Eth denominations. Users are authenticated through their MetaMask wallet which ensures the credibility of every individual player. A smart contract is used to securely transfer and receive cryptocurrency during betting and payouts. These transactions are recorded in the blockchain, so any fraudulent activity will be recorded and caught. The platform also ensures fairness by allowing players to adjust the AI to play at their skill level. The AI component brings a fair and challenging component to the game. Blackjack AI is a platform that focuses on entertainment and education for the players while allowing them to interact with AI players and bet cryptocurrency.

1.3 Motivating User Story

Online games today are full of microtransactions and advertisements. Small developer-run games normally are full of bugs and not fun to use. Now, with a bit of Testnet Sepolia Ethereum and a MetaMask account, anybody can play a fun blackjack game right alongside a real artificial intelligence agent. Motivated users can even create programs that interface directly with our website to play as well!

Blackjack AI is motivated to bring the best user experience to the customers. It prioritizes safety and entertainment over all else. Users who are struggling with a gambling addiction can use Blackjack AI to get the thrill of betting on blackjack while playing with friends, without financial consequences.

2 Implementation Details

2.1 AI Aspects

We used the reinforcement learning algorithm "Q-learning" to train our AI model. We chose this algorithm because it lends itself to AI exploration of game environments, where there is a reward that is easily modeled, as the purpose of Q-learning is to maximize that reward. Due to the nature of reinforcement learning, an agent makes decisions over time in an environment, and those decisions affect its knowledge of the environmental states, which impacts a decision's probability of being chosen. We refer to a Q-learning AI agent as a QAgent from here on. A QAgent's knowledge is stored in a so-called "Q-table", which is a mapping of environment states to reward values, each associated with a decision. Over time, these reward values are adjusted, and the one that becomes maximized is the decision that the agent will learn to choose when provided its associated state. We performed an initial training session of 5,000,000 blackjack games to create an effective Q-table with optimal strategy. This allowed our QAgent to maintain a level of "smartness", which we could control by providing the probability of the QAgent selecting a decision from the Q-table, thus being the probability of the QAgent making the optimal decision. Our QAgent model came from an open-source Jupyter Notebook [1], written by Till Zemann, which can be found in the bibliography.

We provided an AI template for users to upload custom agents that contain a fixed set of methods that are executed within the program. The abstract Agent class acts as the skeleton for all agents used in the Blackjack AI system, including the QAgent. For an agent to interact with the system it must have certain methods defined, so they can be executed and the behavior can be guaranteed by the system. We provided an example of a custom agent that either hits or stands via a "coin flip", or a 50% chance of making either choice.

2.2 Blockchain Aspects

Our blockchain implementation is written and compiled in Solidity and deployed on the Sepolia ETH test network through RemixIDE. It contains private member variables that hold a mapping of addresses to balances, as well as a list of all users that have participated. The constructor saves the address of the sender and declares it as the owner. The contract is initialized as "payable", and can receive ETH/Wei at any time, however, the main payable function that will be used is called "deposit". This function not only accepts ETH/Wei into the contract but also saves the amount deposited to the user's account in the mapping. If the user has a balance of 0, their address is also added to the user list. The cashOut function is used to withdraw money from a user's account. It has a modifier that requires the requested value to cash out to be within the range of their current balance. The function sends the user the requested amount of ETH/Wei and then updates their balance. The changeBalance function is meant to be an administrator-only function, as it directly modifies the user's balance without receiving or sending any currency. This function is used to adjust balances during the betting/winning processes. To safeguard this function, it requires an input of the owner's wallet address to proceed. This function also protects against negative integers by ensuring that if the input value is larger than the user's balance, it sets their balance to 0. The contract also contains two view functions: getBalance and contractBalance. The function getBalance is used to return the balance of the user that is stored in the blockchain, and contractBalance returns the full amount of ETH/Wei that the contract has. The final function is meant to be a development-only function, called evacuateFunds. It loops across all users that have deposited into the contract and sends them any amount of money that they have in the contract. The leftover money is then returned to the owner of the contract. The main use of this function is to allow the deployer and users to not lose their money if a new version of the contract is deployed.

2.3 AI and Blockchain Integration

Both the AI and Blockchain aspects of our project are integrated through the Python backend and JavaScript frontend. The AI's Python implementation was specifically defined in a manner that allows it to be wrapped by a custom-created object called LocalPlayer. This object is used to wrap all of the possible agents that can play the game, from the AI player to the WebUser (which is controlled almost exclusively through the frontend). This wrapper acts as an interface between the backend game logic and the AI agent. This

3 Project Demonstration

To use the Blackjack AI platform, the user first needs to log in with their MetaMask account, as shown in Figure 1. Upon opening the website for the first time, they will be automatically prompted to link their account. After the first login, the user will only ever be prompted to log in if they need to unlock their wallet, as MetaMask saves a list of connected sites. The user has the option to deposit or withdraw Sepolia ETH into and from their balance at will. Next, the user has the option to upload an agent to play alongside them. This agent undergoes a testing process that verifies its validity and if rejected, the user will be notified. A description of how to create an agent that fits this validity is listed in the README. The user can also use a slider to adjust the level of our provided pre-trained AI. The slider runs from 1-100, where 100 is the smartest. The AI's display name on the next screen depends on this "smartness" value. The smartest bot is named 'Lord Blackjack', at 100% smartness. Anywhere below 100 but above 70 grants the AI the name "Master Mind AI". If the smartness is between 40 and 70, the AI will be named "Ninja AI". A smartness value between 10 and 40 warrants the name "NPC AI", and a level below 10 grants the AI the name "Fresh

Off the Compiler AI”.

After logging in, the user is taken to the game’s main page. The buttons ”DEAL”, ”BET”, STAND”, and ”HIT”, allow the user to play the game and bet the Sepolia ETH in their balance. The main user is displayed in the middle of the screen so that their cards are perfectly vertical for easy reading. On the right of the player will always be the site-provided AI. If the user provided an agent file, their agent will be displayed on the left side of the table, with the username ”;User;’s Agent”, where ”User” is the user’s username.

The ”BET” button becomes active at the beginning of each game, as shown in Figure 2. To place a bet, users can adjust the desired amount of Sepolia Ethereum (in Wei) by pressing the ”+” or ”-“ buttons. The buttons will not allow a user to bet higher than their balance, nor can they bet negatively. When the user is satisfied with their amount, they must press the ”BET” button to confirm the bet. This will trigger a MetaMask transaction that removes the bet amount from the user’s wallet. This transaction will pop up again if rejected, as the user must confirm the transaction to continue.

Once the bet is confirmed, the ”DEAL” button becomes active. Two cards are dealt to each player and the dealer. The players can see their cards, or they can click on the cards to flip them over and hide them from view. According to the rules of blackjack, only one of the dealer’s cards is visible.

After the cards are dealt, the ”STAND” and ”HIT” buttons become active. The player can hit when they want to be dealt another card from the deck, and they can stand when they are done. This portion of the game will be automatically ended on a user stand, or if the user hits and subsequently busts.

As seen in Figure 3, the user will then be prompted with another MetaMask transaction to confirm the results of the game. If the player wins, they will be paid back twice their bet. If the player was dealt blackjack, they will be paid back thrice their bet. If the player and the dealer push, they will be paid back their bet, and if the player loses, they will not be paid anything. As with the bet transaction, this transaction will continue to pop up if the user rejects it, as they cannot proceed without updating their balance.

The player can play as many rounds as they want, or until they run out of funds in their balance. In that case, the player should use the ”EXIT” button to return to the login page and deposit more funds into their account.

As a note about the persistent transactions: Users can avoid these transactions by refreshing the page, as the site will not remember the game that they were in. However, this either means that the user was able to cancel their bet and must bet again, or they have canceled a payout and lost some of their balance. Because these two cases do not encourage dodging transactions, we did not feel the need to implement data that persists beyond refreshing.

4 Evaluation

Our program’s main goal was to have an AI agent that played blackjack optimally, that a user could play alongside. We plotted the results of the AI training over 5,000,000 games, as seen in Figure 5. Over time, the rewards were maximized, and the training error curve flattened out until our AI was playing mostly optimal blackjack. We also plotted the strategy generated by the AI, which can be seen in Figures 6 and 7. We can observe that the AI’s strategy was more conservative when was not holding a usable ace, as it was more likely to bust, versus its strategy when it was holding a usable ace, as it could afford to be more risky. To test the blockchain aspect of our code, we made a spreadsheet that aimed to test corner cases of the functionality of our contract. We deployed a test version of the contract that was used exclusively to conduct the tests. These tests are reflected in Figure , and the spreadsheet itself is linked in the resources section. These tests ensured that the fail-safes that we intended to program for did work. This included ensuring that users could not withdraw more money than what was in their balance, ensuring that users could not call the balance modification function easily, and ensuring that only the owner could remove all of the money from the contract. This document also tested the standard functionality of the contract, including deposits, standard withdrawals, adjusting balances after pretend wins and losses, and removing all funds.

As for the frontend and backend that ran and interacted with both the AI agent(s) and the smart contract, we covered as many edge cases as we could consider. As we were developing the website we ran into a few errors here and there and were able to correct them, but we also thought of some exploits that we needed to protect against. First off, both the backend and the frontend do not allow the execution of functions that are not expected. For example, a user cannot deal themselves cards, hit, stand, and receive payment

without first placing a bet and confirming it in the frontend and subsequently the backend. The backend game logic uses a state enumeration to proceed to the next state, only allowing functions to run if the active state matches the required state of the function. Likewise, the frontend uses boolean values to keep track of if a button should be able to be pressed if a button was pressed at the wrong time, the function would simply exit. We took this one step further and disabled buttons that were not meant to be clicked, fully preventing any wrong interaction. As discussed in the demo section, we also ensured that users cannot dodge signing their MetaMask transactions without fully refreshing the webpage, ensuring that they never gain currency from this exploit. We also do not allow users to access the blackjack page unless they go through the login page, as any refresh attempts redirect the user to the login page. Finally, we created a redirect from any parent directory URL to the main login screen, allowing users to enter a generic URL and be immediately directed to the login page.

5 Conclusion

5.1 Summary

Our project allows users to play blackjack games alongside artificial intelligence agents and experience a fully functional blockchain that acts as a bank to keep track of balances. Users can upload custom Agent files, which play alongside them at the blackjack table.

5.2 Ethical Use

Our project is purely simulative, so there are no concerns regarding real financial systems or currencies. As our program allows users to upload custom artificial intelligence agents, we made sure to have some protective measures against system abuse put in place. The uploaded agent code is checked against a series of verifiers to ensure that the specified methods are defined within the class. This allows our program to guarantee that no errors will be thrown when executing the custom agent code, which would throttle our system. Another protective measure that was implemented is the fact that only certain predefined class methods are defined within the program. This guarantees that our system can interpret the behavior of the agent. The use of the agent code is wrapped in a try-except block, so if any methods fail we ignore the custom class, which protects against users who do not follow the template and function behavior exactly. If we were to deploy this system, the end user would not know anything about the system's code structure, apart from the custom artificial intelligence agent template. We would also implement a timeout checker, by wrapping each custom-declared function in a timing decorator, which would allow our system to be protected against code that either runs for too long or is maliciously intending to slow down the system. Lastly, the MetaMask integration itself has preventative measures against abuse, which we use within our system. The maximum amount a user can bet is their MetaMask Testnet Sepolia Ethereum balance. This verification is performed by MetaMask and the "Testnet" so no abuse of this fact is possible, as this case is handled externally.

5.3 Limitations

Blackjack AI has some limitations that we need to address before releasing the game to the public. The blackjack game does not implement doubling down or splitting. This is mainly due to the limitations of the artificial intelligence model that we used. The Q-Table is only capable of generating hit and stand as decisions. Furthermore, the site is not correctly configured to host multiple users at one time, it can only serve a single user. This would need to be ratified before full deployment.

5.4 Future Work

We do not plan to maintain or update the project. Ultimately there are plenty of ways for anyone to play Blackjack, and while we have a cool concept, we do not believe it is worth the time, effort, and money to deploy to a live website and continue to maintain the project.

5.5 Project Resources

Github

Contract address and code are also stored in the Github repository.

Blockchain Testing Document

Demo

References

- [1] *Gymnasium documentation*. Solving Blackjack with Q-Learning - Gymnasium Documentation. (n.d.). https://gymnasium.farama.org/tutorials/training_agents/blackjack_tutorial/

6 Figures

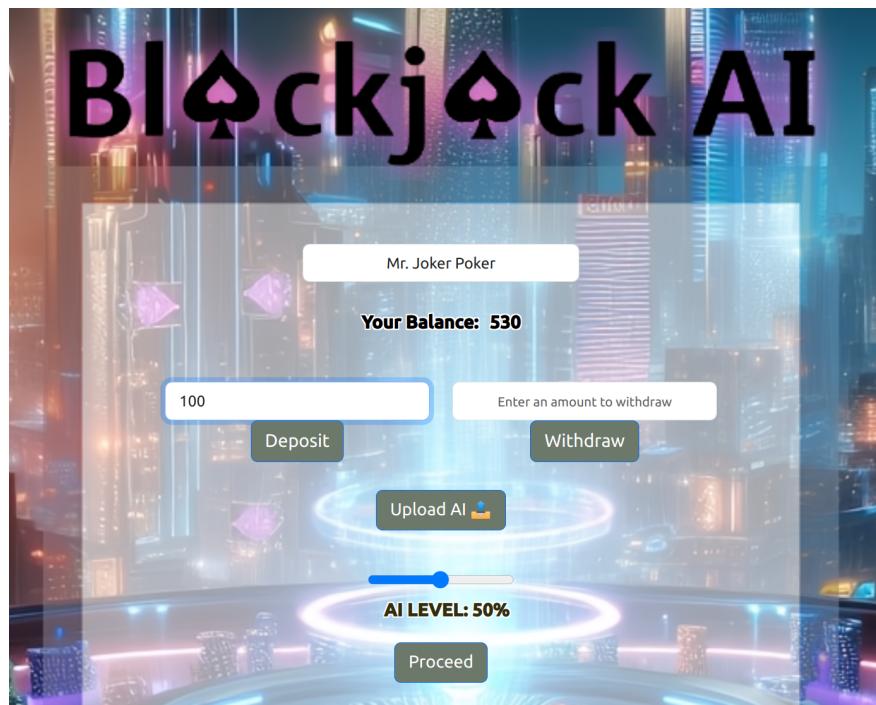


Figure 1: The Blackjack AI Login page

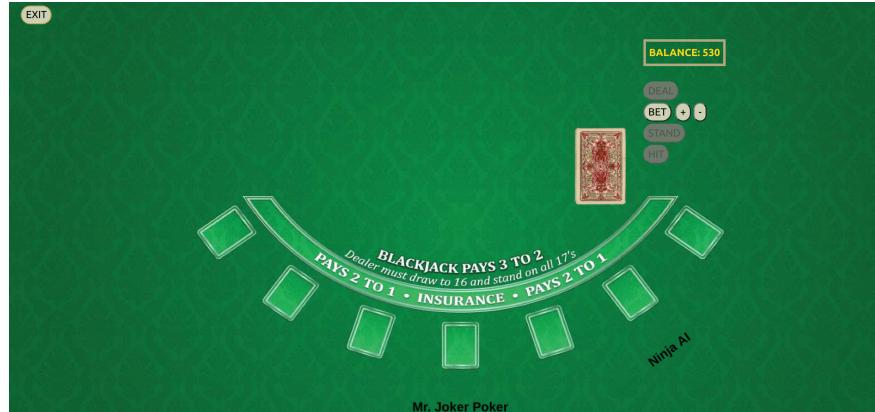


Figure 2: Game Page: The "BET" button is active at the beginning the game

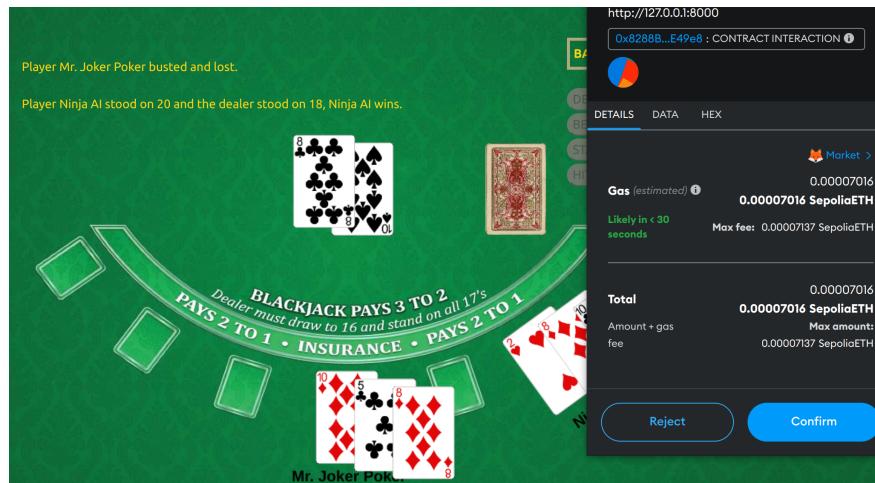


Figure 3: Transaction at the end of each round

Text Content Address:	Expected Answer	Actual Answer	TX hash (if applicable)	Image no./hash
Can I accept currency on deployment?	Yes	Yes	0x3d320e237ab75947e5447308666129883772320x92d8e941	Figure 1
Does user deposit money?	Yes	Yes	0x7494e82b30c1827203b7130648954571989f8c2e72052500a480975	Figure 2
Does this account automatically add their balance? Is it queried?	Yes	Yes	N/A	Figure 3
Does the contract automatically add the change?	Yes	Yes	N/A	
Is a user able to change a balance in their wallet, set the user's balance overflow?	No	No	0xb03833a5d37b9d9e537a3ff1fffb1a1e1d92667d71901dd287	Figure 4
Can a user change their balance without the owner's address as input?	No	No	0xb2c32413c1e3d875c0d1e1e173e881310ff70fc0ac873d7497409	Figure 5
Can a user change their balance with the owner's address as input?	Yes	Yes	0x2654d72052722957265eb0383151565e5a533bd8763283a27	Figure 6
Can a user withdraw money whenever they want?	Yes	Yes	0xa2ca119703b42e11163048876a399151209e176991107715e5a8	Figure 7
Can a user successfully withdraw money whenever they want?	Yes	Yes	0x3520235203bda4e5af95c2e58b5d0f8eaee20b489347c3d0e6d	Figure 8
Can someone who is not the owner use the execute funds button?	No	No	0x6565116473e0ed1a85cc525beee563a11545e1e0300e77221aa35326f	Figure 9
When the owner uses execute funds, is every participant successfully repaid the correct amount?	Yes	Yes	0x6565116473e0ed1a85cc525beee563a11545e1e0300e77221aa35326f	Figure 10
Likewise, is the owner returned the remainder of the money?	Yes	Yes	0x6565116473e0ed1a85cc525beee563a11545e1e0300e77221aa35326f	Figure 11

Figure 4: Blockchain testing document

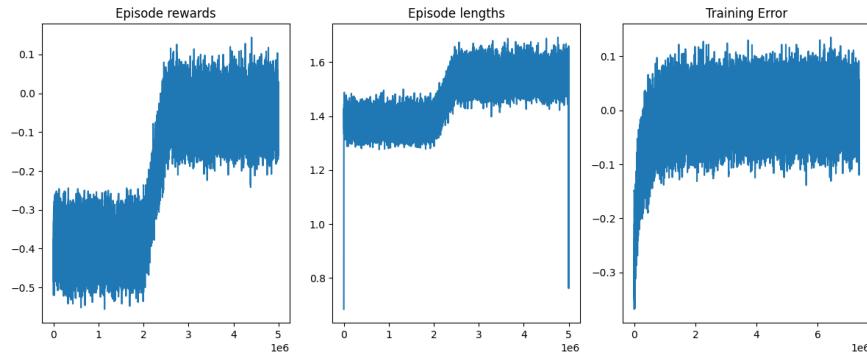


Figure 5: Q-Learning Training Results

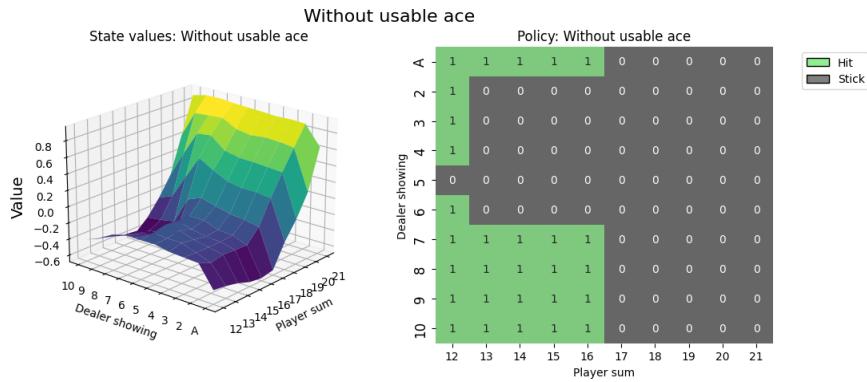


Figure 6: AI Strategy: No Usable Ace

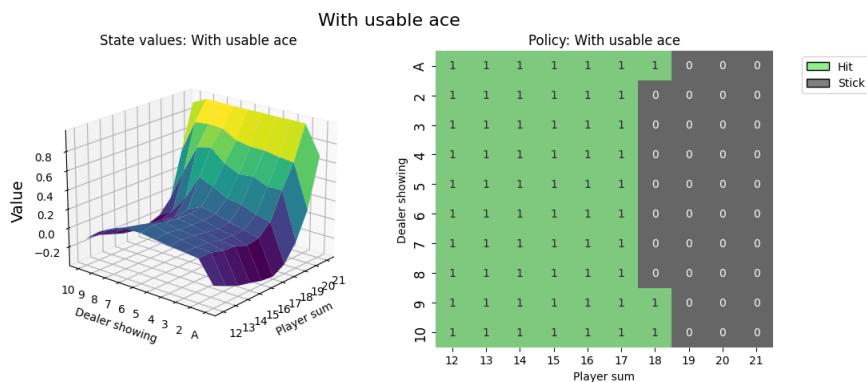


Figure 7: AI Strategy: With Usable Ace